

Ollama Cloud Models: How They Change Prose Drafting

Ollama's **Cloud Models** feature allows you to run frontier-scale models (up to 671B+ parameters) directly through the familiar Ollama CLI/API interface without needing high-end enterprise GPUs locally. The models are offloaded to Ollama's cloud infrastructure while maintaining seamless compatibility with local writing tools like SillyTavern, Obsidian plugins, or custom scripts.

Top Ollama Cloud Models for Prose:

- **deepseek-v3.1:671b-cloud:** The flagship open model for prose. Rivals top proprietary models in subtext and character voice without heavy refusal behaviors.
- **gpt-oss:120b-cloud:** Offers frontier-level structural capacity for complex, long-form chapters and detailed scene beats.
- **kimi-k2:1t-cloud:** Built for massive context and fluid descriptive writing; ideal for multi-chapter continuity.
- **glm-4.6:cloud:** Highly nuanced tone control and elegant sentence structuring.
- **gpt-oss:20b-cloud:** Snappy, ultra-fast responses for quick line editing, synonym brainstorming, and punchy dialogue generation.



Top 20 AI models for writing prose, featuring proprietary leaders, specialized fiction engines, open-weights flagships, and the latest **Ollama Cloud-hosted models**.

Top 20 AI Models for Prose Writing

Rank Model Type / Host Key Prose Strengths & Best Use Cases **1 Anthropic Claude Opus** Proprietary API The **Prose Champion**. Gold-standard sentence rhythm, rich subtext, deep character voice retention. Best for literary fiction & climactic chapters. **2 Anthropic Claude Sonnet** Proprietary API Unmatched balance of prose quality, rapid iteration speed, and minimal AI cliché tropes. The default workhorse for most fiction authors. **3 Sudowrite Muse** Specialized Platform Purpose-built for fiction; fine-tuned on published novels. Exceptional for narrative description and beats-to-prose expansion. **4 DeepSeek-v3.1 (671B Cloud)** **Ollama Cloud** / API High-reasoning open architecture with deep literary nuance, sharp dialogue, and impressive subtext handling. **5 OpenAI GPT-5 / GPT-4o** Proprietary API Outstanding structural logic, macro-level plotting, high-stakes scene pacing, and tight thriller framing. **6 Kimi K2 / K3 (1T Cloud)** **Ollama Cloud** / API Massive context handling with strong descriptive flow, long-form narrative consistency, and expressive imagery. **7 Google Gemini Pro** Proprietary API Multi-million token context window; perfect for loading massive series bibles, world-building lore, and historical research. **8 GLM-4.6 (Cloud)** **Ollama Cloud** / API Excellent multi-turn pacing, poetic phrasing, and clean handling of complex tone shifts in narrative arcs. **9 Mistral Large** Proprietary API / Open Sophisticated European styling, balanced sentence variation, and strong multilingual

literary prose. **10 NovelAI (Kayra / Clio)** Specialized Platform Purpose-built for fiction with zero restrictive content filters; granular style control for genre fiction (romance, dark fantasy). **11 GPT-OSS 120B (Cloud)** **Ollama Cloud** Large-scale open-weights model offloaded via Ollama Cloud; offers unconstrained creative control and rich world-building output. **12 Qwen 2.5 / 3 (72B / Max)** Open / Ollama Cloud / API Rich sensory imagery, strong scene setting, fantasy/sci-fi lore generation, and budget-friendly API access. **13 Llama 3.3 70B Local** / Hosted The open-source benchmark for local setups; highly responsive to custom system prompts and fiction fine-tunes. **14 Hermes 3 70B (Nous)** Local / Hosted API Unrestricted roleplay, deep immersive narrative writing, and raw character interactions without corporate voice hedging. **15 DeepSeek-R1 (Reasoning)** Local / Cloud API Outstanding for internal monologues, complex character psychology, and intricate plot-point troubleshooting. **16 Cohere Command R+** Proprietary API Clean, highly objective non-fiction prose; excellent for narrative essays, whitepapers, and journalistic profiles. **17 GPT-OSS 20B (Cloud)** **Ollama Cloud** Lightweight, low-latency cloud model via Ollama; ideal for rapid dialogue drafting and quick scene iterations. **18 Yi-Lightning** / **Yi-34B** Open / Hosted Smooth conversational pacing, expressive descriptive prose, and strong emotional cadence in dialogue. **19 Mixtral 8x22B Local** / Open Hosted High-speed multi-agent story drafting, quick dialogue generation, and solid narrative momentum. **20 XAI Grok** Proprietary API Sharp, casual, and punchy tone; excels at humor, satire, snappy banter, and unvarnished social commentary.