

path-scurry

Extremely high performant utility for building tools that read the file system, minimizing filesystem and path string munging operations to the greatest degree possible.

Ugh, yet another file traversal thing on npm?

Yes. None of the existing ones gave me exactly what I wanted.

Well what is it you wanted?

While working on [glob](#), I found that I needed a module to very efficiently manage the traversal over a folder tree, such that:

1. No `readdir()` or `stat()` would ever be called on the same file or directory more than one time.
2. No `readdir()` calls would be made if we can be reasonably sure that the path is not a directory. (Ie, a previous `readdir()` or `stat()` covered the path, and `ent.isDirectory()` is `false`.)

Benchmarks can be run by executing `npm run bench`.
 for convenience.

PathSecurity provides ample opportunities to trade performance
 good performance, as well as several options to get extremely
 fast, in the cases it's optimized for, if used properly.

hope that this module doesn't blaze anyone. But it does go very
 fast, in the cases it's optimized for, if used properly.

JavaScript people throw around the word "blazing" a lot. I
 hope that this module doesn't blaze anyone. But it does go very
 fast, in the cases it's optimized for, if used properly.

PERFORMANCE

having to parse or join or do any error-prone string munging.
 consumer can return canonical platform-specific paths without
 paths and drive letters and wrongway slashes, so that the
 7. Handle all the weird aspects of Windows paths, like UNC
 directory with a huge number of entries.

6. Do not blow up the JS heap allocation if operating on a
 and go get it seamlessly when requested.

5. It's more important to prevent excess syscalls than to be up to
 date, but it should be smart enough to know what it *doesn't* know,
 many cases where an API like node's `f5` is preferable.

4. The API is not limited to use as a stream/iterator/etc. There are
 don't have to ever check again).

we find that they *don't* exist, store that information as well, so we
 it has to track "provisional" child nodes that may not exist (and if
 string-parsing/munging operations are be minimized. This means
 3. `path.resolve()`, `dirname()`, `basename()`, and other

As is always the case, doing more means going slower, doing less means going faster, and there are trade offs between speed and memory usage.

PathScurry makes heavy use of [LRUCache](#) to efficiently cache whatever it can, and Path objects remain in the graph for the lifetime of the walker, so repeated calls with a single PathScurry object will be extremely fast. However, adding items to a cold cache means “doing more”, so in those cases, we pay a price. Nothing is free, but every effort has been made to reduce costs wherever possible.

Also, note that a “cache as long as possible” approach means that changes to the filesystem may not be reflected in the results of repeated PathScurry operations.

For resolving string paths, PathScurry ranges from 5-50 times faster than `path.resolve` on repeated resolutions, but around 100 to 1000 times *slower* on the first resolution. If your program is spending a lot of time resolving the *same* paths repeatedly (like, thousands or millions of times), then this can be beneficial. But both implementations are pretty fast, and speeding up an infrequent operation from 4 μ s to 400ns is not going to move the needle on your app’s performance.

For walking file system directory trees, a lot depends on how often a given PathScurry object will be used, and also on the walk method used.

With default settings on a folder tree of 100,000 items, consisting of around a 10-to-1 ratio of normal files to directories, PathScurry performs comparably to [@nodelib/fs.walk](#), which is the fastest and most reliable file system walker I could find. As far

On my machine, that is about 1000-1200 completed walks per second for async or stream walks, and around 500-600 walks per second for sync walks, and around 4x faster for stream walks than sync walks. In the warm cache state, PathScurity's performance increases around 4x for sync for await iteration, 10-15x faster for streams and syncrous for of iteration, and anywhere from 30x to 80x faster for the rest.

# operations / ms	New PathScurity object Reuse	PathScurity object	sync stream: 492.718 15028.343 async walk: 1095.648 32706.395 sync walk: 527.632 46129.772 async sync: 1288.821 5045.510 sync sync: 498.496 17920.746
A hand-rolled walk calling entry.readDir()	The cold cache state is still limited by the costs of file system operations, but with a warm cache, the only bottleneck is CPU operations, but with a warm cache, the only bottleneck is CPU speed and VM optimizations. Of course, in that case, some care must be taken to ensure that you don't lose performance as a result as I can tell, it's almost impossible to go much faster in a Node.js program, just based on how fast you can push syscalls out to the fs thread pool.		
# walk 100,000 fs entries, 10/1 file/dir ratio	In the warm cache state, PathScurity's performance increases around 4x for sync for await iteration, 10-15x faster for streams and syncrous for of iteration, and anywhere from 30x to 80x faster for the rest.		
30x to 80x faster for the rest.	30x to 80x faster for the rest.		
second synchronous.	second synchronous.		
On my machine, that is about 1000-1200 completed walks per second for async or stream walks, and around 500-600 walks per second for sync walks, and around 4x faster for stream walks than sync walks. In the warm cache state, PathScurity's performance increases around 4x for sync for await iteration, 10-15x faster for streams and syncrous for of iteration, and anywhere from 30x to 80x faster for the rest.			
30x to 80x faster for the rest.	30x to 80x faster for the rest.		
ratio	ratio		
# operations / ms	# operations / ms		
PathScurity object	PathScurity object		
sync stream: 492.718 15028.343 async walk: 1095.648 32706.395 sync walk: 527.632 46129.772 async sync: 1288.821 5045.510 sync sync: 498.496 17920.746	sync stream: 492.718 15028.343 async walk: 1095.648 32706.395 sync walk: 527.632 46129.772 async sync: 1288.821 5045.510 sync sync: 498.496 17920.746		
A hand-rolled walk calling entry.readDir()	A hand-rolled walk calling entry.readDir()		
The cold cache state is still limited by the costs of file system operations, but with a warm cache, the only bottleneck is CPU speed and VM optimizations. Of course, in that case, some care must be taken to ensure that you don't lose performance as a result as I can tell, it's almost impossible to go much faster in a Node.js program, just based on how fast you can push syscalls out to the fs thread pool.	The cold cache state is still limited by the costs of file system operations, but with a warm cache, the only bottleneck is CPU speed and VM optimizations. Of course, in that case, some care must be taken to ensure that you don't lose performance as a result as I can tell, it's almost impossible to go much faster in a Node.js program, just based on how fast you can push syscalls out to the fs thread pool.		

If the nearest common ancestor is the root, then an absolute path is returned.

path.relativePosix(): string

Return the relative path from the PathWalker cwd to the supplied path string or entry, using / path separators.

If the nearest common ancestor is the root, then an absolute path is returned.

On posix platforms (ie, all platforms except Windows), this is identical to pw.relative(path).

On Windows systems, it returns the resulting string as a /-delimited path. If an absolute path is returned (because the target does not share a common ancestor with pw.cwd), then a full absolute UNC path will be returned. Ie, instead of 'C:\\foo\\bar', it would return //?/C:/foo/bar.

async path.readdir()

Return an array of Path objects found by reading the associated path entry.

If path is not a directory, or if any error occurs, returns [], and marks all children as provisional and non-existent.

path.readdirSync()

Synchronous path.readdir()

of silly mistakes, like calling readdir() on entries that you know are not directories.

# manual recursive iteration functions	
	cold cache warm cache
async:	1164.901 17923.320
cb:	1101.127 40999.344
zalgo:	1082.240 66689.936
sync:	526.935 87097.591

In this case, the speed improves by around 10-20x in the async case, 40x in the case of using entry.readdirCB with protections against synchronous callbacks, and 50-100x with callback deferrals disabled, and *several hundred times faster* for synchronous iteration.

If you can think of a case that is not covered in these benchmarks, or an implementation that performs significantly better than PathScurry, please [let me know](#).

USAGE

```
// hybrid module, load with either
// method
import { PathScurry, Path } from 'path-
scurry'
// or:
const { PathScurry, Path } =
require('path-scurry')
```

```

// very simple example, say we want to
// delete all the .DS_Store files in a
// given path
// note that the API is very similar to
// native walk with fs.readdir()
import { unlink } from 'fs/promises'
// easy way, iterate over the directory
const pw = new PathScrubber(process.cwd())
for await (const entry of pw) {
  if (entry.isFile() && entry.name ===
    '.DS_Store') {
    await unlink(entry.fullPath())
  }
}

// here it is as a manual recursive
const walk = async (entry: Path) => {
  // readdir doesn't throw on non-
  // directories, it just doesn't
  // return any entries, to save stack
  // trace costs.
  const promises: Promise<any> = []
  for (const child of await
    pw.readdir(entry)) {
    // items are returned in arbitrary
    // unsorted order
    if (child.isDirectory()) {
      // each child is a Path object
      // we can add it to our stack
      promises.push(walk(child))
    }
  }
  // here we join all the promises
  // and do the thing
  const result = await Promise.all(promises)
  return result
}

// here it is as a manual recursive
// method
method
{
  // here it is as a manual recursive
  // method
  const walk = async (entry: Path) => {
    // readdir doesn't throw on non-
    // directories, it just doesn't
    // return any entries, to save stack
    // trace costs.
    const promises: Promise<any> = []
    for (const child of await
      pw.readdir(entry)) {
      // each child is a Path object
      // we can add it to our stack
      promises.push(walk(child))
    }
    // here we join all the promises
    // and do the thing
    const result = await Promise.all(promises)
    return result
  }
}

```

<p><code>path.relative()</code></p> <p>Return the relative path from the PathWalker cwd to the supplied path string or entry.</p>	<p><code>path.resolve(p: string)</code></p> <p>Return a Path object associated with the provided path string as resolved from the current Path object.</p>
<p><code>path.resolve()</code></p> <p>When the path is known to not exist, returns true if the path's type is unknown. Always returns true</p>	<p><code>path.resolve(p: string)</code></p> <p>When the path is known to not exist, returns true if the path's type is unknown. Always returns true</p>
<p><code>path.isUnknown()</code></p> <p>Same as the identical <code>fs.Dirent.isX()</code> methods.</p>	<p><code>path.isUnknown()</code></p> <p>Same as the identical <code>fs.Dirent.isX()</code> methods.</p>
<p><code>path.isDirectory(), path.isDirectory(), etc.</code></p> <p>On windows, this will return a fully resolved absolute UNC path using / separators. Eg, instead of 'C:\\foo\\bar', it will return '//?/C:/foo/bar'.</p>	<p><code>path.isDirectory()</code>, <code>path.isDirectory()</code>, etc.</p> <p>On posix systems, this is identical to <code>path.fullpath()</code>.</p>
<p>The fully resolved path to the entry, using / separators.</p>	<p>The fully resolved path to the entry, using / separators.</p>
<p><code>path.fullpath()</code></p> <p>The fully resolved path to the entry.</p>	<p><code>path.fullpathposix()</code></p> <p>The fully resolved path to the entry.</p>

`path.isNamed(name: string): boolean`

Return true if the path is a match for the given path name. This handles case sensitivity and unicode normalization.

Note: even on case-sensitive systems, it is **not** safe to test the equality of the `.name` property to determine whether a given pathname matches, due to unicode normalization mismatches.

Always use this method instead of testing the `path.name` property directly.

`path.isCWD`

Set to true if this Path object is the current working directory of the `PathScurry` collection that contains it.

`path.getType()`

Returns the type of the Path object, '`File`', '`Directory`', etc.

`path.isType(t: type)`

Returns true if `is{t}()` returns true.

For example, `path.isType('Directory')` is equivalent to `path.isDirectory()`.

`path.depth()`

Return the depth of the Path entry within the directory tree.

Root paths have a depth of 0.

```
// could also do pw.resolve(entry,
// child.name),
// just like fs.readdir walking,
// but .fullpath is
// a *slightly* more efficient
// shorthand.
promises.push(unlink(child.fullpath()))
} else if (child.isDirectory()) {
  promises.push(walk(child))
}
}
return Promise.all(promises)
}

walk(pw.cwd).then(() => {
  console.log('all .DS_Store files
    removed')
})

const pw2 = new PathScurry('/a/b/c') //
  pw2.cwd is the Path for /a/b/c
const relativeDir = pw2.cwd.resolve('../
  x') // Path entry for '/a/b/x'
const relative2 = pw2.cwd.resolve('/a/b/
  d../x') // same path, same
  entry
assert.equal(relativeDir, relative2)
```

API

[Full TypeDoc API](#)

Class Path implements fs.Dirent

Object representing a given path on the file system, which may

1. *Leucosia* 2. *Leucosia* 3. *Leucosia* 4. *Leucosia*

in use. They differ in the separators used to split and join path

strings, and the handling of root paths.

In Pathway3, implementations paths are split using either

in use based on the drives and UNC paths encountered. UNC paths such as //?/C:/ that identify a drive letter, will be treated as an alias for the same root entry as their associated drive letter. //, or \\\, and joined using \\\, and multiple roots may be (in this case 'C:\\').

path.name

Important: always test the path name againts any test string using the `isNamed` method, and not by directly comparing this string. Otherwise, unicode path strings that the system sees as identical will not be properly treated as the same path, leading to incorrect behavior and possible security issues.

- **Name of this file system entry.**
Important: always test the path name against any test string using the `isNamed` method, and not by directly comparing this string. Otherwise, unicode path strings that the system sees as identical will not be properly treated as the same path, leading to incorrect behavior and possible security issues.

Interface Paths Configuration Options

There are platform-specific classes exported, but for the most part, the default PathScurry and Path exports are what you most likely need, unless you are testing behavior for other platforms.

Intended public API is documented here, but the full documentation does include internal types, which should not be accessed directly.

The type of the options argument passed to the

- nocase: Boolean indicating that file names should be compared case-insensitively. Defaults to true on darwin and false elsewhere.

Warning Performing case-insensitive matching on a case-sensitive file system may negatively impact performance.

- **childренCacheSize**: Number of child entries to cache, in order to speed up `resolve()` and `readDir()` calls. Defaults

Setting it to a higher value will run the risk of JS heap allocation errors on large directory trees. Setting it to 256 or smaller will significantly reduce the construction time and data consumption overhead, but with the downside of operations being slower on large directory trees. Setting it to 0 will mean that to 100, 100+ (cc, 100+).

```
async pw.lstat(entry = pw.cwd)
```

Call `fs.lstat` on the supplied string or Path object, and fill in as much information as possible, returning the updated Path object.

Returns `undefined` if the entry does not exist, or if any error is encountered.

Note that some `Stats` data (such as `ino`, `dev`, and `mode`) will not be supplied. For those things, you'll need to call `fs.lstat` yourself.

```
pw.lstatSync(entry = pw.cwd)
```

Synchronous `pw.lstat()`

```
pw.realpath(entry = pw.cwd, opts = { withFileTypes: false })
```

Call `fs.realpath` on the supplied string or Path object, and return the realpath if available.

Returns `undefined` if any error occurs.

May be called as `pw.realpath({ withFileTypes: boolean })` to run on `pw.cwd`.

```
pw.realpathSync(entry = pw.cwd, opts = { withFileTypes: false })
```

Synchronous `pw.realpath()`

effectively no operations are cached, and this module will be roughly the same speed as `fs` for file system operations, and *much* slower than `path.resolve()` for repeated path resolution.

- `fs` An object that will be used to override the default `fs` methods. Any methods that are not overridden will use Node's built-in implementations.
 - `lstatSync`
 - `readdir (callback withFileTypes Dirent variant, used for readdirCB and most walks)`
 - `readdirSync`
 - `readlinkSync`
 - `realpathSync`
 - `promises: Object containing the following async methods:`
 - `lstat`
 - `readdir (Dirent variant only)`
 - `readlink`
 - `realpath`

Interface WalkOptions

The options object that may be passed to all walk methods.

- `withFileTypes`: Boolean, default true. Indicates that Path objects should be returned. Set to `false` to get string paths instead.
- `follow`: Boolean, default false. Attempt to read directory entries from symbolic links. Otherwise, only actual directories are

```

Syncronous pw.readlink()
{
    pw.readlinkSync(link = pw.cwd, opts = {
        withFileTypes: false
    })
}

Note that TypeScript return types will only be inferred properly from static analysis if the withFileTypes option is omitted, or a constant true or false value.

Returns undefined if any error occurs (for example, if the argument is not a symbolic link), or a Path object if the argument is a symbolic link, or a string otherwise.

Can be called as pw.readlink({ withFileTypes: boolean }) as well.

```

Call fs.readlink on the supplied string or Path object, and return the result.

```

async pw.readlink(link = pw.cwd, opts = {
    withFileTypes: false
})

```

Synchronous pw.readdir()

```

{
    pw.readdirSync(dir = pw.cwd, opts = {
        withFileTypes: true
    })
}

Note that TypeScript return types will only be inferred properly from static analysis if the withFileTypes option is omitted, or a constant true or false value.

Returns [] if no entries are found, or if any error occurs.

Can be called as pw.readdir({ withFileTypes: boolean }) as well.

```

- **filter: Function (entry: Path) => boolean.**
Sets this imposes a slight performance penalty, because readdir must be called on all symbolic links encountered, in traversed directory will never be followed.
traversed. Regardless of this setting, a given target path will only ever be walked once, meaning that a symbolic link to a previously traversed directory will be avoided in order to avoid infinite cycles.
- **filterer: Function (entry: Path) => boolean.**
Setting this imposes a slight performance penalty, because readdir this imposes a slight performance penalty, because filterer does not pass the filter, though it will prevent the traversed if they do not pass the filter, although it will prevent the filterer from being included in the results. By default, if no filter is provided, then all entries are included in the results.
- **walkFilter: Function (entry: Path) => boolean.**
If provided, will prevent the traversal of any directory (or in the case of follow: true symbolic links to directories) for which the function returns false. This will not prevent the directories themselves from being included in the result set. Use filter for that.
- **follow: boolean.**
Note that TypeScript return types will only be inferred properly from static analysis if the withFileTypes option is omitted, or a constant true or false value.
- **pathsCurry: Class.**
The main interface. Defaults to an appropriate class based on the current platform.

Class PathsCurry

```
pw.relativePosix(path: string | Path):  
string
```

Return the relative path from the PathWalker cwd to the supplied path string or entry, using / path separators.

If the nearest common ancestor is the root, then an absolute path is returned.

On posix platforms (ie, all platforms except Windows), this is identical to `pw.relative(path)`.

On Windows systems, it returns the resulting string as a /-delimited path. If an absolute path is returned (because the target does not share a common ancestor with `pw.cwd`), then a full absolute UNC path will be returned. Ie, instead of 'C:\foo\bar, it would return //?/C:/foo/bar.

```
pw.basename(path: string | Path): string
```

Return the basename of the provided string or Path.

```
pw.dirname(path: string | Path): string
```

Return the parent directory of the supplied string or Path.

```
async pw.readdir(dir = pw.cwd, opts =  
{ withFileTypes: true })
```

Read the directory and resolve to an array of strings if `withFileTypes` is explicitly set to `false` or Path objects otherwise.

Use `PathScurryWin32`, `PathScurryDarwin`, or `PathScurryPosix` if implementation-specific behavior is desired.

All walk methods may be called with a `WalkOptions` argument to walk over the object's current working directory with the supplied options.

```
async pw.walk(entry?: string | Path |  
WalkOptions, opts?: WalkOptions)
```

Walk the directory tree according to the options provided, resolving to an array of all entries found.

```
pw.walkSync(entry?: string | Path |  
WalkOptions, opts?: WalkOptions)
```

Walk the directory tree according to the options provided, returning an array of all entries found.

```
pw.iterate(entry?: string | Path |  
WalkOptions, opts?: WalkOptions)
```

Iterate over the directory asynchronously, for use with `for await of`. This is also the default async iterator method.

```
pw.iterateSync(entry?: string | Path |  
WalkOptions, opts?: WalkOptions)
```

Iterate over the directory synchronously, for use with `for of`. This is also the default sync iterator method.

`pw.stream(entry?: string | Path | WalkOptions, opts?: WalkOptions)`

Return a [Minipass](#) stream that emits each entry or path string in the walk. Results are made available asynchronously.

`pw.streamSync(entry?: string | Path | WalkOptions, opts?: WalkOptions)`

Return a [Minipass](#) stream that emits each entry or path string in the walk. Results are made available synchronously, meaning that the walk will complete in a single tick if the stream is fully consumed.

`pw.cwd()`

Path object representing the current working directory for the [PathSync](#).

`pw.chdir(path: string)`

Set the new effective current working directory for the security object, so that `path.relative()` and `path.relativePosix()` return values relative to the new cwd path.

`pw.depth(path?: Path | string): number`

Return the depth of the specified path (or the `PathSync cwd`) within the directory tree.

Root entries have a depth of 0.

`pw.stream(entry?: string | Path | WalkOptions, opts?: WalkOptions)`

Return a [Minipass](#) stream that emits each entry or path string in the walk. Results are made available asynchronously.

`pw.streamSync(entry?: string | Path | WalkOptions, opts?: WalkOptions)`

Return a [Minipass](#) stream that emits each entry or path string in the walk. Results are made available synchronously, meaning that the walk will complete in a single tick if the stream is fully consumed.

`pw.cwd()`

Path object representing the current working directory for the [PathSync](#).

`pw.chdir(path: string)`

Set the new effective current working directory for the security object, so that `path.relative()` and `path.relativePosix()` return values relative to the new cwd path.

`pw.depth(path?: Path | string): number`

Return the depth of the specified path (or the `PathSync cwd`) within the directory tree.

Root entries have a depth of 0.

If the nearest common ancestor is the root, then an absolute path is returned.

`pw.relative(path?: string | Path): string`

Return the relative path from the `PathWalker` cwd to the supplied path string or entry.

`pw.relativePosix(path?: string | Path): string`

Return the relative path from the `PathWalker` cwd to the new effective current working directory for the security object, so that `path.relative()` and `path.relativePosix()` return values relative to the new cwd path.

`pw.resolve(paths: string[]): string`

Caching `path.resolve()`, but always using posix style separators. Ie, instead of `'C:\\foo\\bar'`, it would return `//?/C:/foo/bar`.

`pw.resolveSync(paths: string[]): string`

On Windows, it returns the full absolute UNC path using '/' separators (ie, everywhere except Windows).

`pw.resolvePaths(paths: string[]): string[]`

This is identical to `pw.resolve('../paths')` on posix paths.

`pw.resolvePosix(paths: string[]): string[]`

Caching `path.resolve()`, but always using posix style paths.

`pw.resolveSync(paths: string[]): string[]`

Takes a single string argument, not multiple.

`pw.resolve(paths: string[]): string`

To get a `Path` object resolved from the `PathSync`, use `pw.cwd().resolve(path)`. Note that `Path.resolve` only builds out the cached `Path` entries.

`pw.resolveSync(paths: string[]): string`

Significantly faster than `path.resolve()` if called repeatedly with the same paths. Significantly slower otherwise, as it builds out the cached `Path` entries.

`pw.resolve(paths: string[]): string`

Caching `path.resolve()`.